

Project Pravega

Storage Reimagined for a Streaming World



技术创新的浪潮接踵而来， 继续搬砖还是奋起直追？

云数据

AI

区块链

架构优化

高效运维

CTO技术选型

微服务

新开源框架

会议：2018年12月07-08日 培训：2018年12月09-10日

地址：北京·国际会议中心



Market Drivers



Massive
Data Growth



Emergence of
Real-Time Apps




Monetize Data
with Analytics



Rapid Dissemination
of Data to Apps



Data Velocity
and Variety

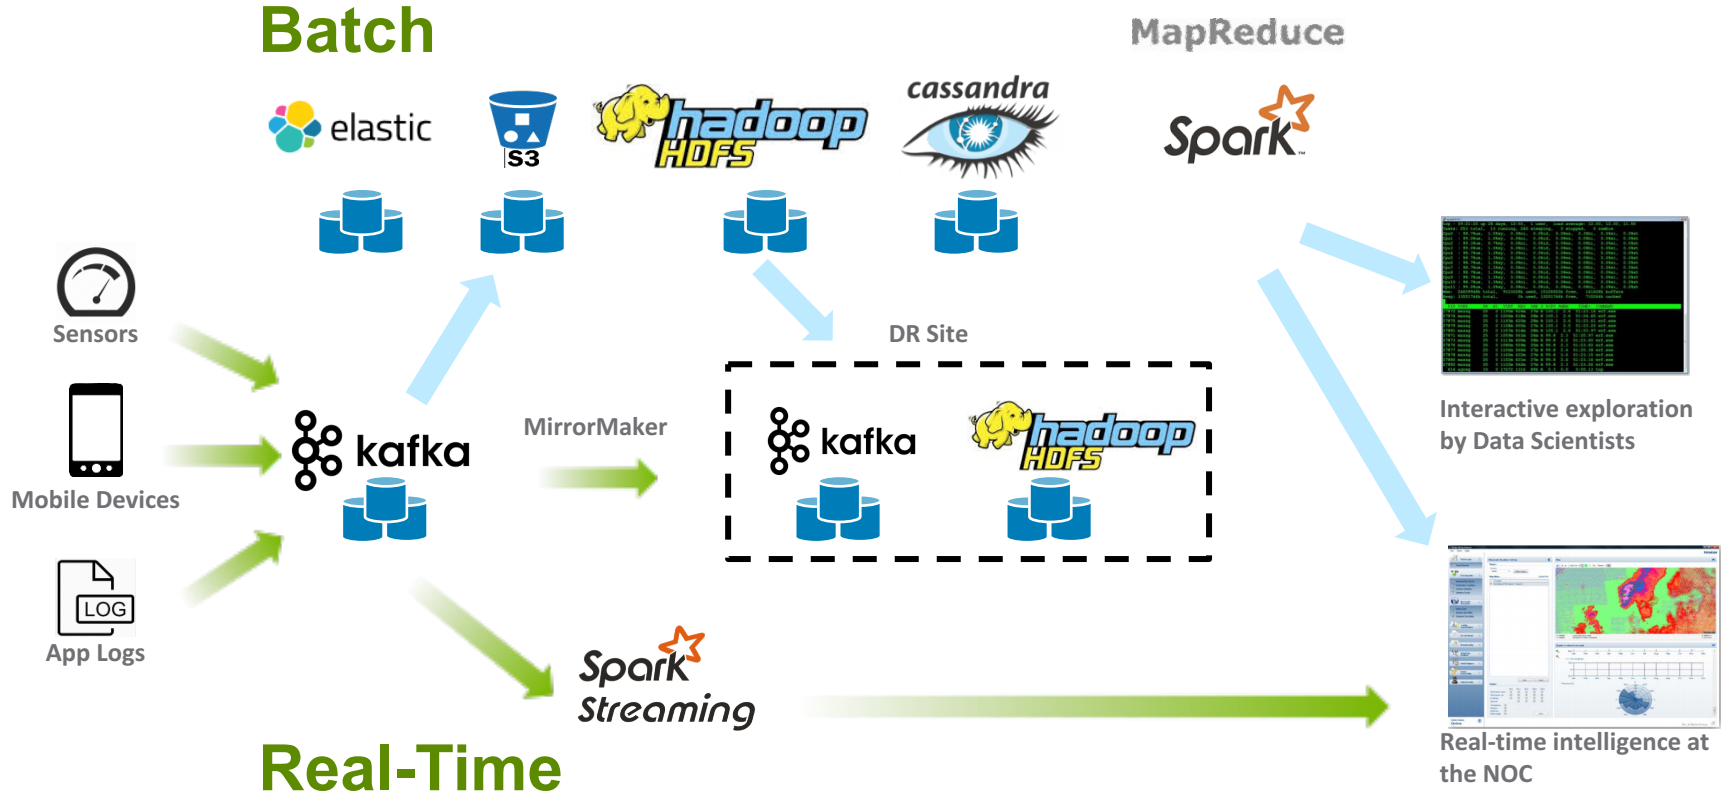


Infrastructure
Commoditization
and Scale-Out



Open Source
Community

Today's "Accidental Architecture"



A New Architecture Emerges: Streaming

- A new class of **streaming systems** is emerging to address the accidental architecture's problems and enable new applications not possible before
- Some of the unique characteristics of streaming applications
 - Treat data as **continuous** and **infinite**
 - Compute **correct results** in **real-time** with stateful, exactly-once processing
- These systems are applicable for real-time applications, batch applications, and interactive applications
- Web-scale companies (Google, Twitter) are beginning to demonstrate the disruptive value of streaming systems
- What are the **implications for storage** in a streaming world?


Let's Rewind A Bit: The Importance of Log Storage

Traditional Apps/Middleware

Streaming Apps/Middleware


BLOCKS

- Structured Data
- Relational DBs



FILES

- Unstructured Data
- Pub/Sub
- NoSQL DBs




OBJECTS

- Unstructured Data
- Internet Friendly (REST)
- Scale over Semantics
- Geo



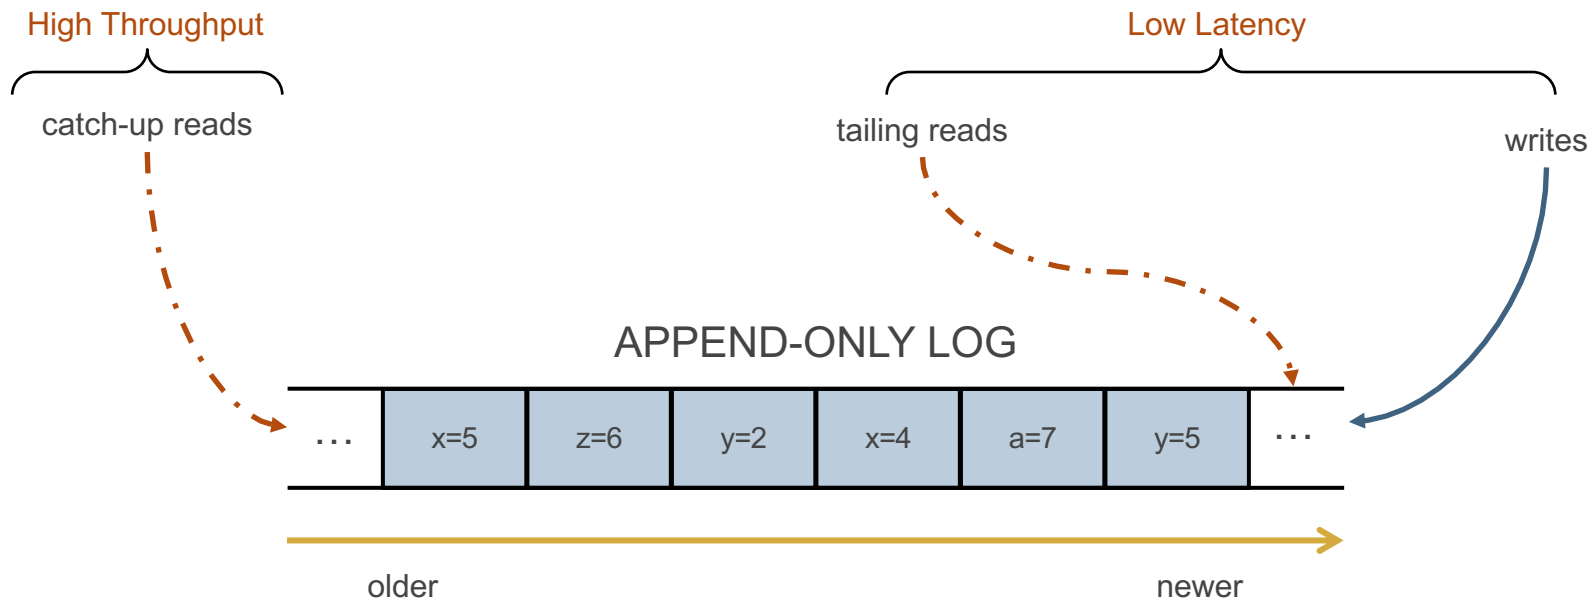
LOGS

- Append-only
- Low-latency
- Tail Read/Write

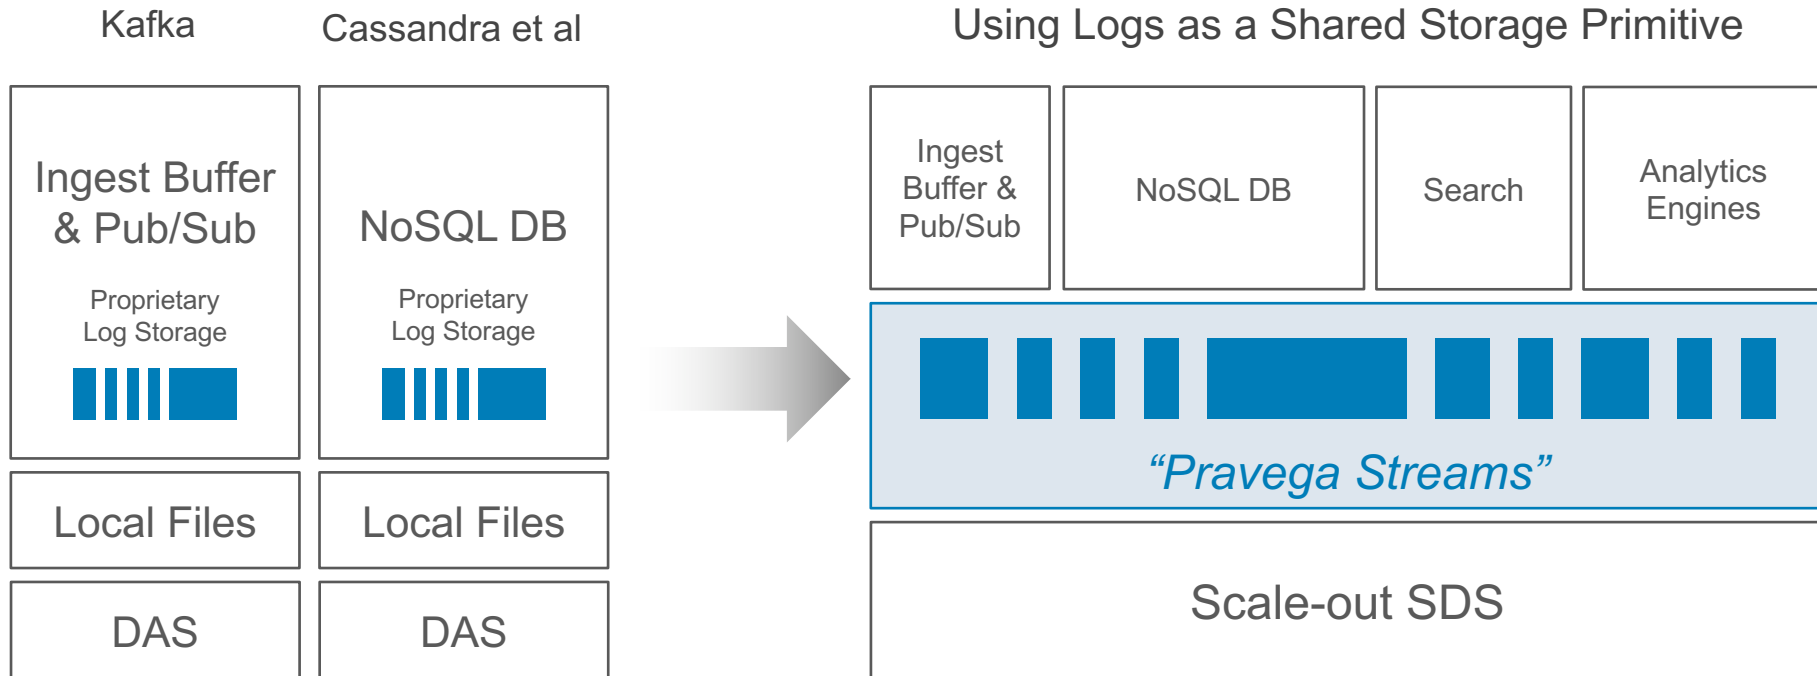


The Importance of Log Storage

The Fundamental Data Structure for Scale-out Distributed Systems



Our Goal: Refactor the “Accidental Storage Stack”



Introducing Pravega Streams

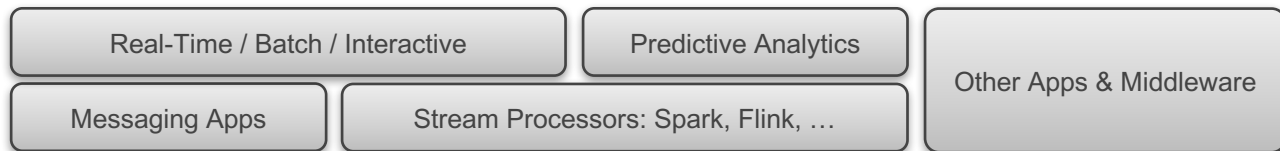
A New Log Primitive Designed Specifically For Streaming Architectures

- Pravega is an open source distributed storage service offering a new storage abstraction called a **stream**
- A stream is the foundation for building reliable streaming systems: a **high-performance, durable, elastic**, and **infinite** append-only log with **strict ordering** and **consistency**
- A stream is as **lightweight as a file** – you can create millions of them in a single cluster
- Streams greatly **simplify** the development and operation of a variety of distributed systems: messaging, databases, analytic engines, search engines, and so on

Pravega Architecture Goals

- All data is durable
 - Data is replicated and persisted to disk before being acknowledged
- Strict ordering guarantees and exactly once semantics
 - Across both tail and catch-up reads
 - Client tracks read offset, Producers use transactions
- Lightweight, elastic, infinite, high performance
 - Support tens of millions of streams
 - Dynamic partitioning of streams based on load and throughput SLO
 - Size is not bounded by the capacity of a single node
 - Low (<10ms) latency writes; throughput bounded by network bandwidth
 - Read pattern (e.g. many catch-up reads) doesn't affect write performance

Architecture



Stream
Abstraction



Cache
(Rocks)

Pravega Streaming Service

Cloud Scale Storage
(Isilon or ECS)

- *High-Throughput*
- *High-Scale, Low-Cost*

Auto-Tiering

Low-Latency Storage

Apache Bookkeeper

Streaming Storage System

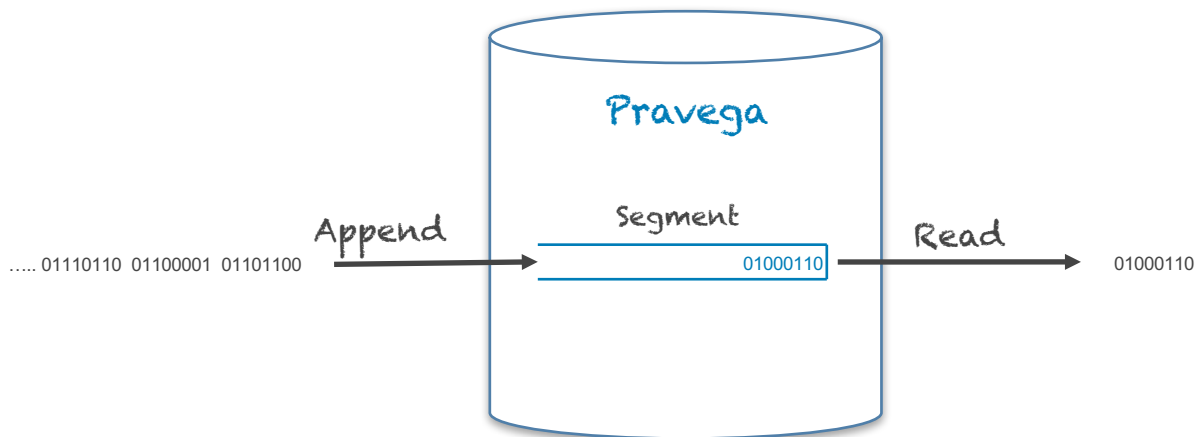
Pravega Design Innovations

1. Zero-Touch Dynamic Scaling
 - Automatically scale read/write parallelism based on load and SLO
 - No service interruptions
 - No manual reconfiguration of clients
 - No manual reconfiguration of service resources
2. Smart Workload Distribution
 - No need to over-provision servers for peak load
3. I/O Path Isolation
 - For tail writes
 - For tail reads
 - For catch-up reads
4. Tiering for "Infinite Streams"
5. Transactions For "Exactly Once"

Pravega Fundamentals


Segments

- Base storage primitive is a segment
- A segment is an append-only sequence of bytes
- Writes durably persisted before acknowledgement



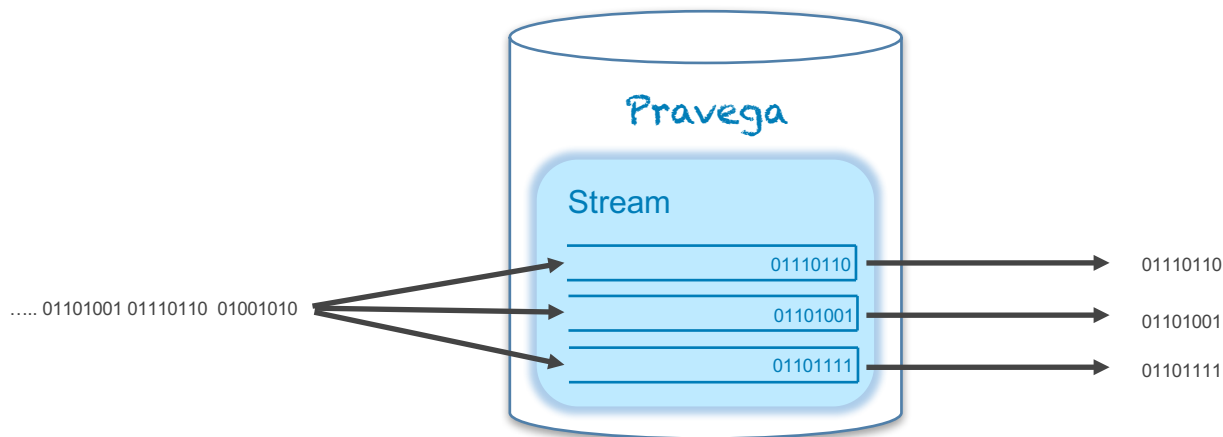
Streams

- A stream is composed of one or more segments
- Routing key determines the target segment for a stream write
- Write order preserved by routing key; consistent tail and catch-up reads



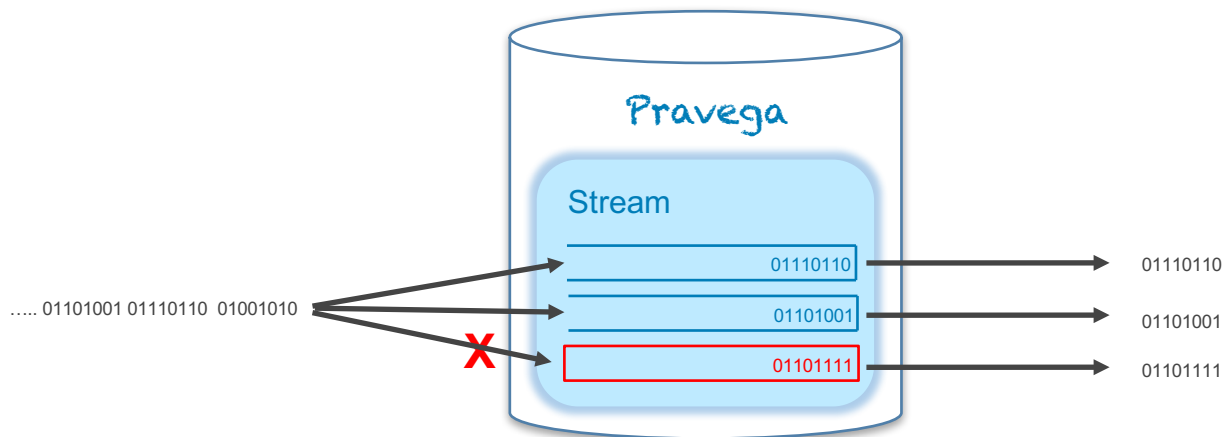
Streams

- There are no architectural limits on the number of streams or segments
- Each segment can live in a different server
- System is not limited in any way by the capacity of a single server

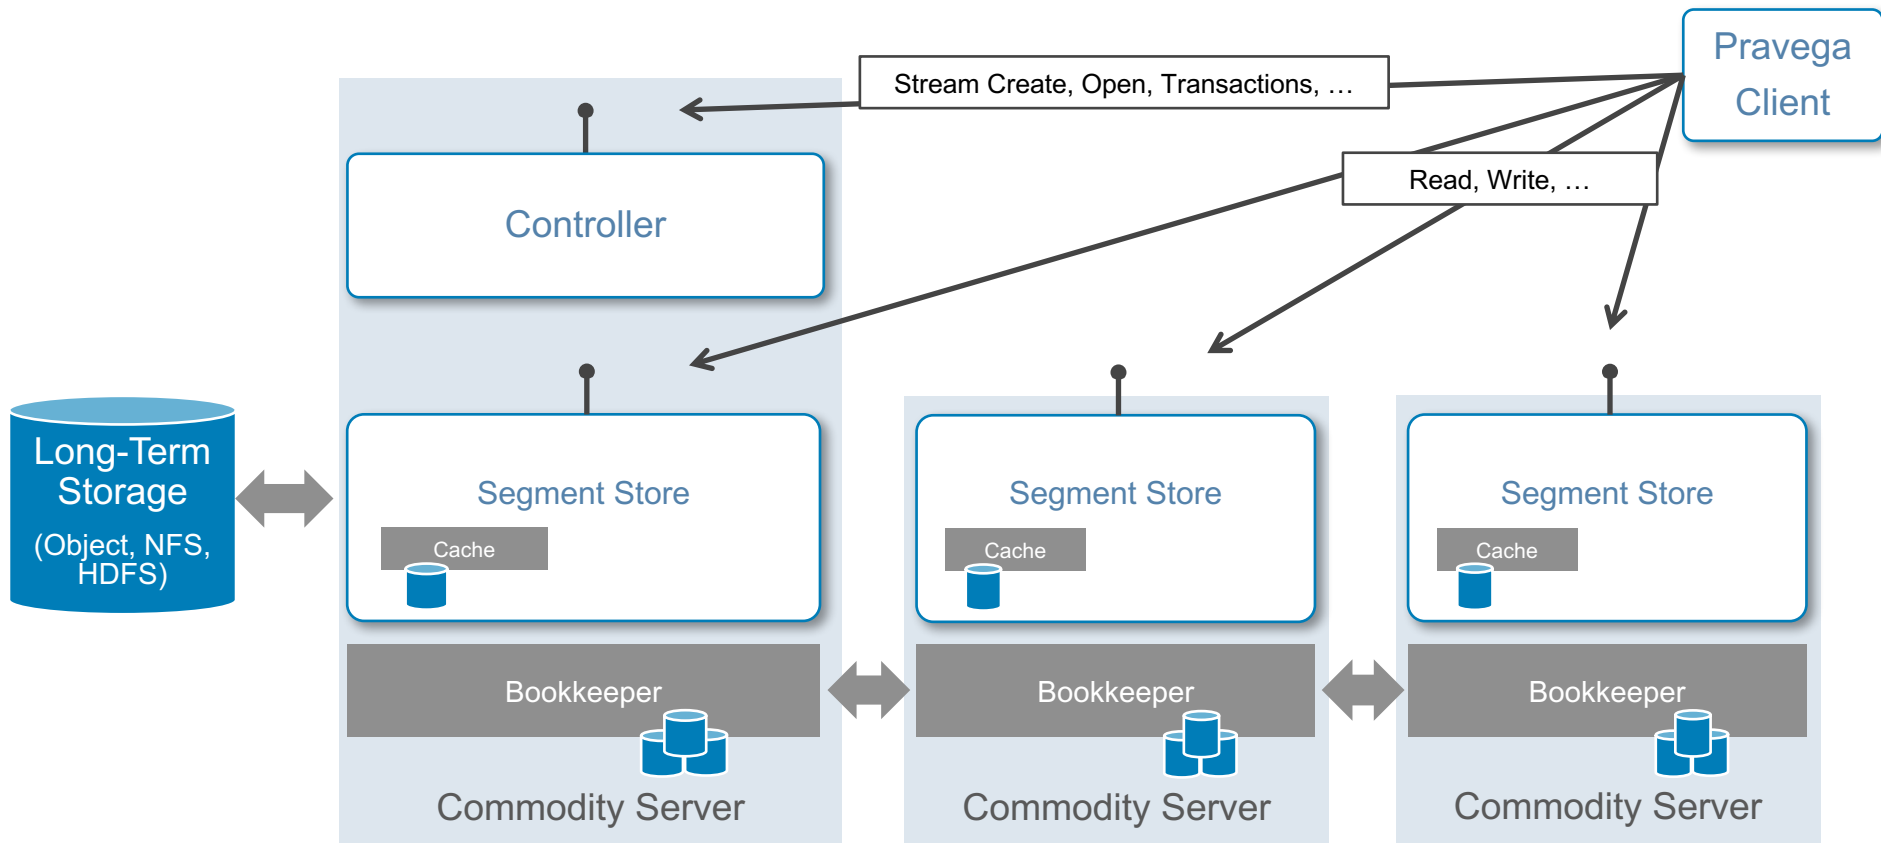


Segment Sealing

- A segment may be sealed
- A sealed segment cannot be appended to any more
- Basis for advanced features such as stream elasticity and transactions



Pravega System Architecture




Beyond the Fundamentals

Stream Elasticity, Unbounded Streams, Transactions, Exactly Once


Stream Elasticity

- Data arrival volume increases – more parallelism needed!

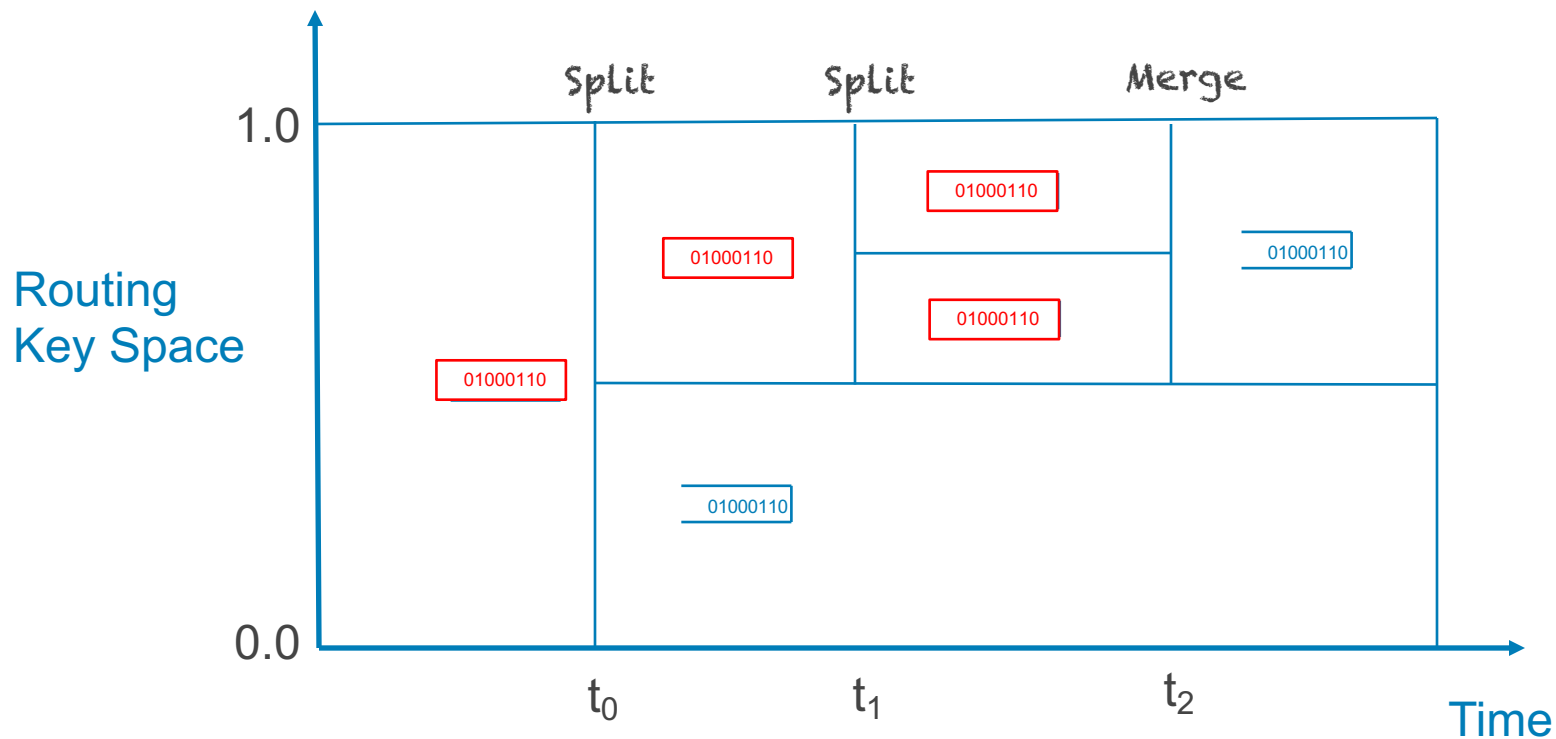


Stream Elasticity

- Seal original segment
- Replace with two new ones!
- New segments may be distributed throughout the cluster balancing load




Stream Elasticity




Stream Elasticity

Key ranges are dynamically assigned to segments




Stream Elasticity

Segment Heat Map

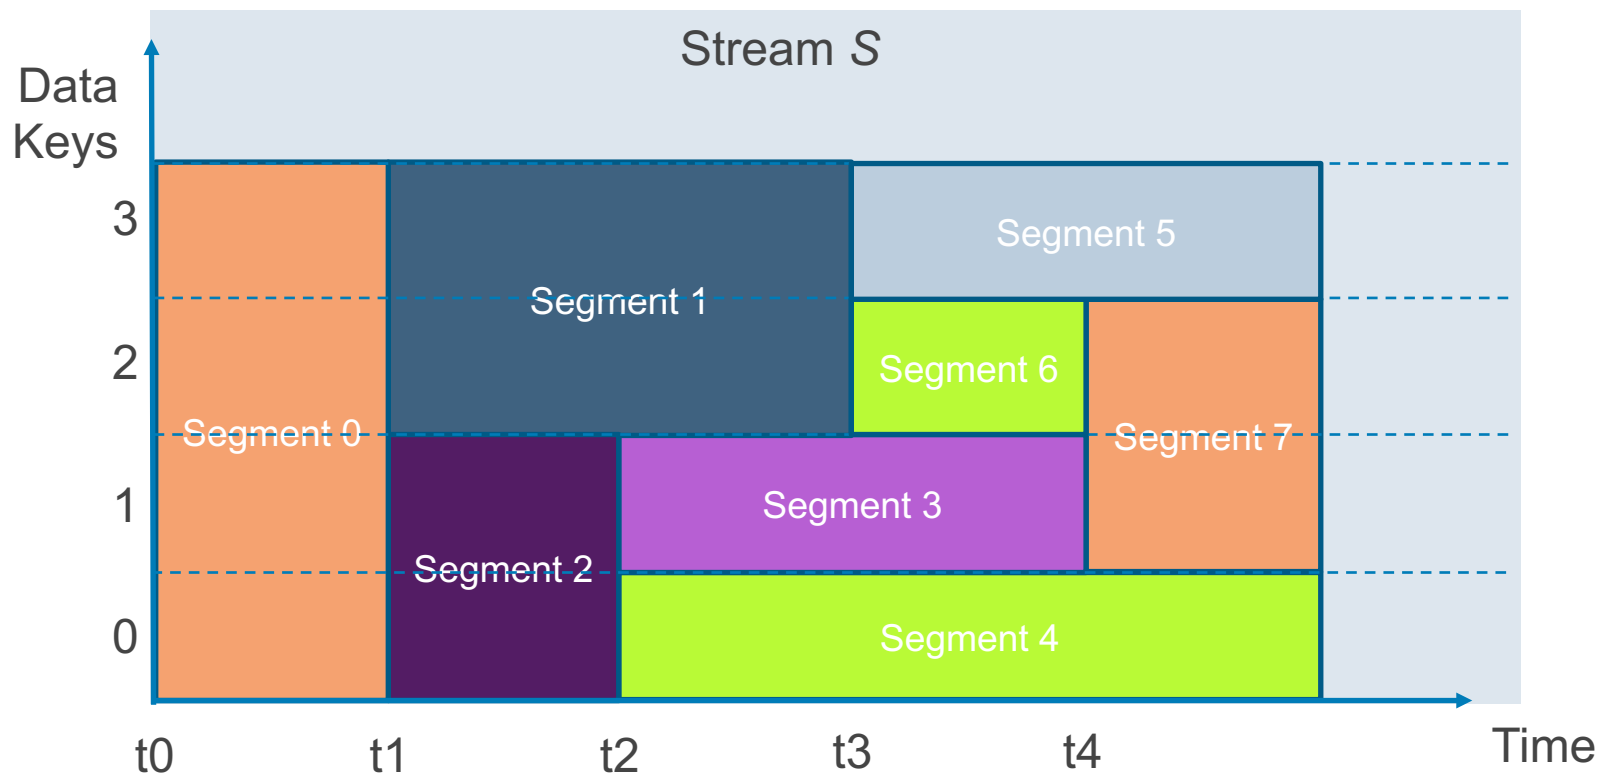


Stream Elasticity

Segment Heat Map

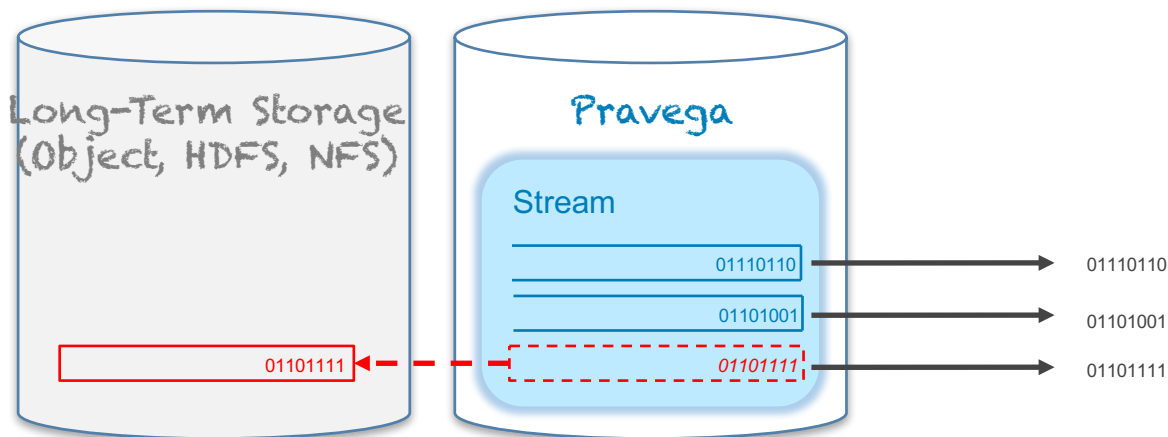


Zero-Touch Scaling: Segment Splitting & Merging







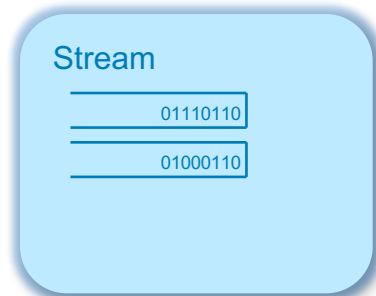
Unbounded Streams

- Segments are automatically tiered to long-term storage
- Data in tiered segments is transparently accessible for catch-up reads
- Preserves stream abstraction while lowering storage costs for older data




Exactly Once

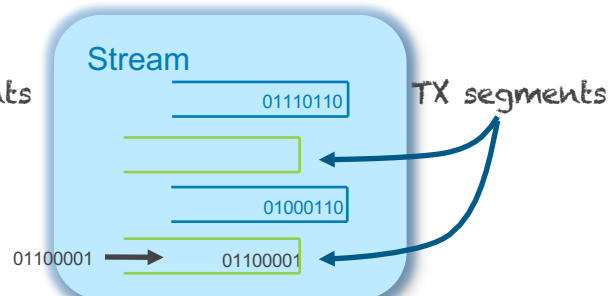




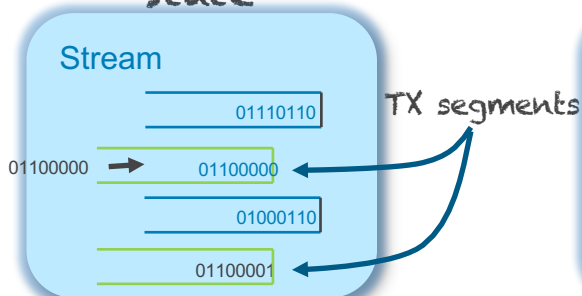
Initial
State




Begin TX




Write to
TX



Write to
TX




Upon
commit



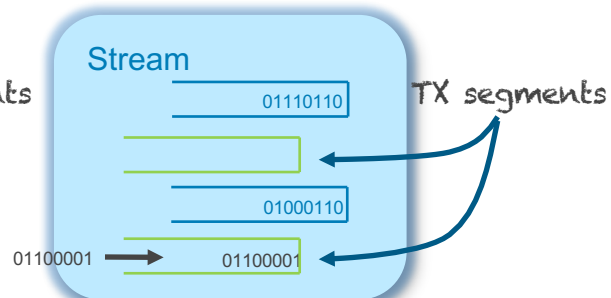
Upon
commit



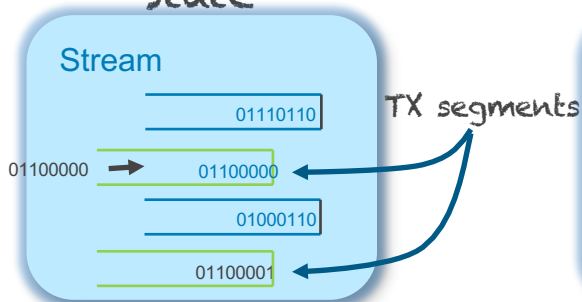
Initial
State




Begin TX



Write to
TX

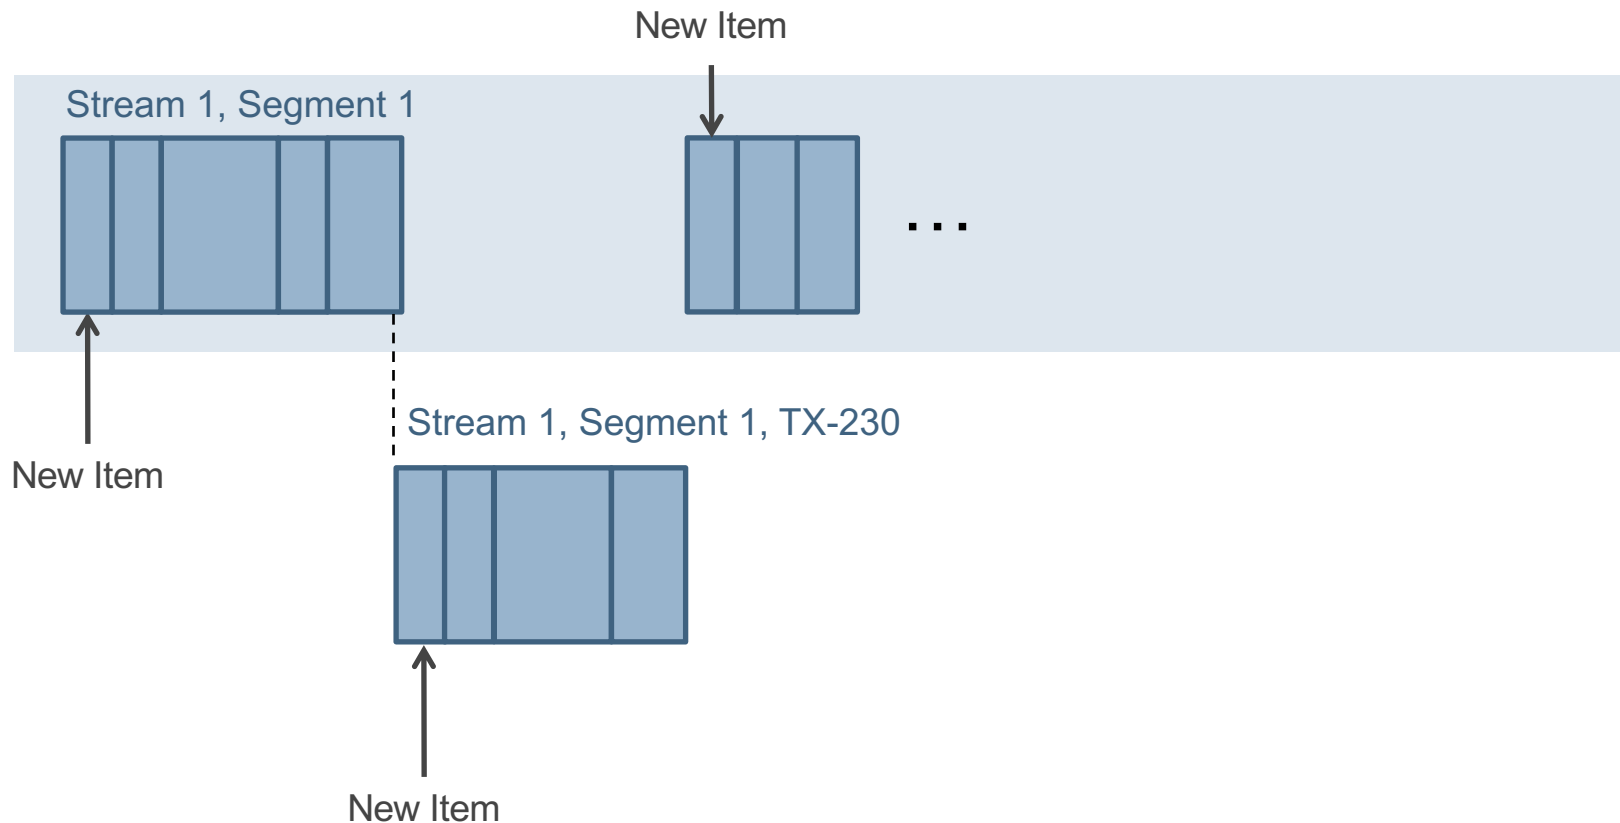


Write to
TX

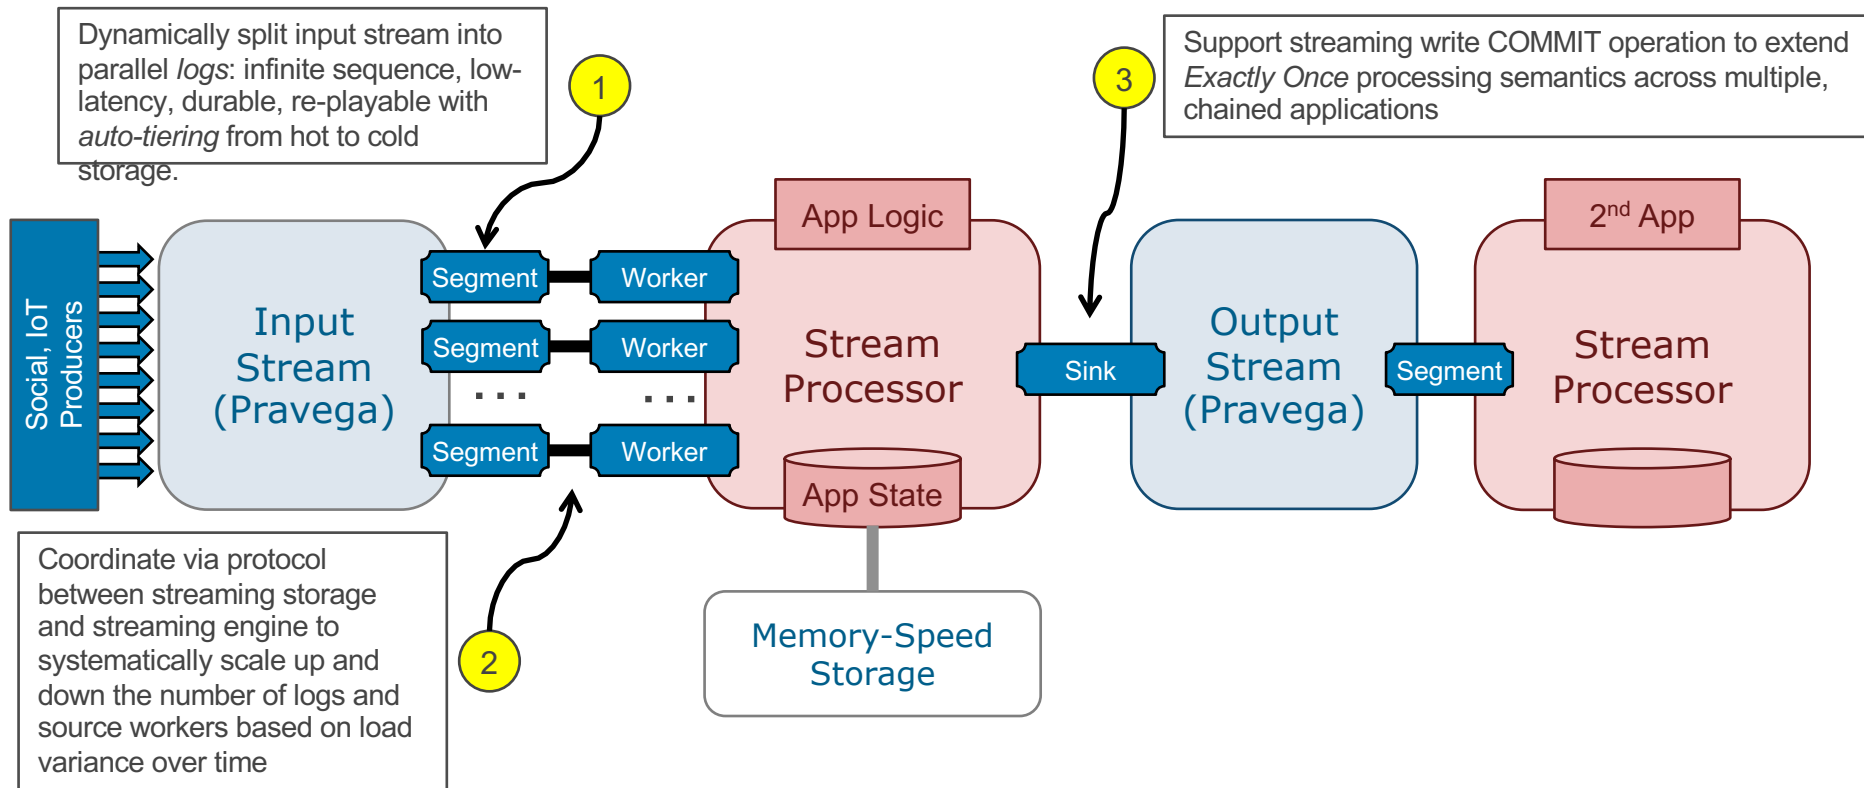


Upon
abort

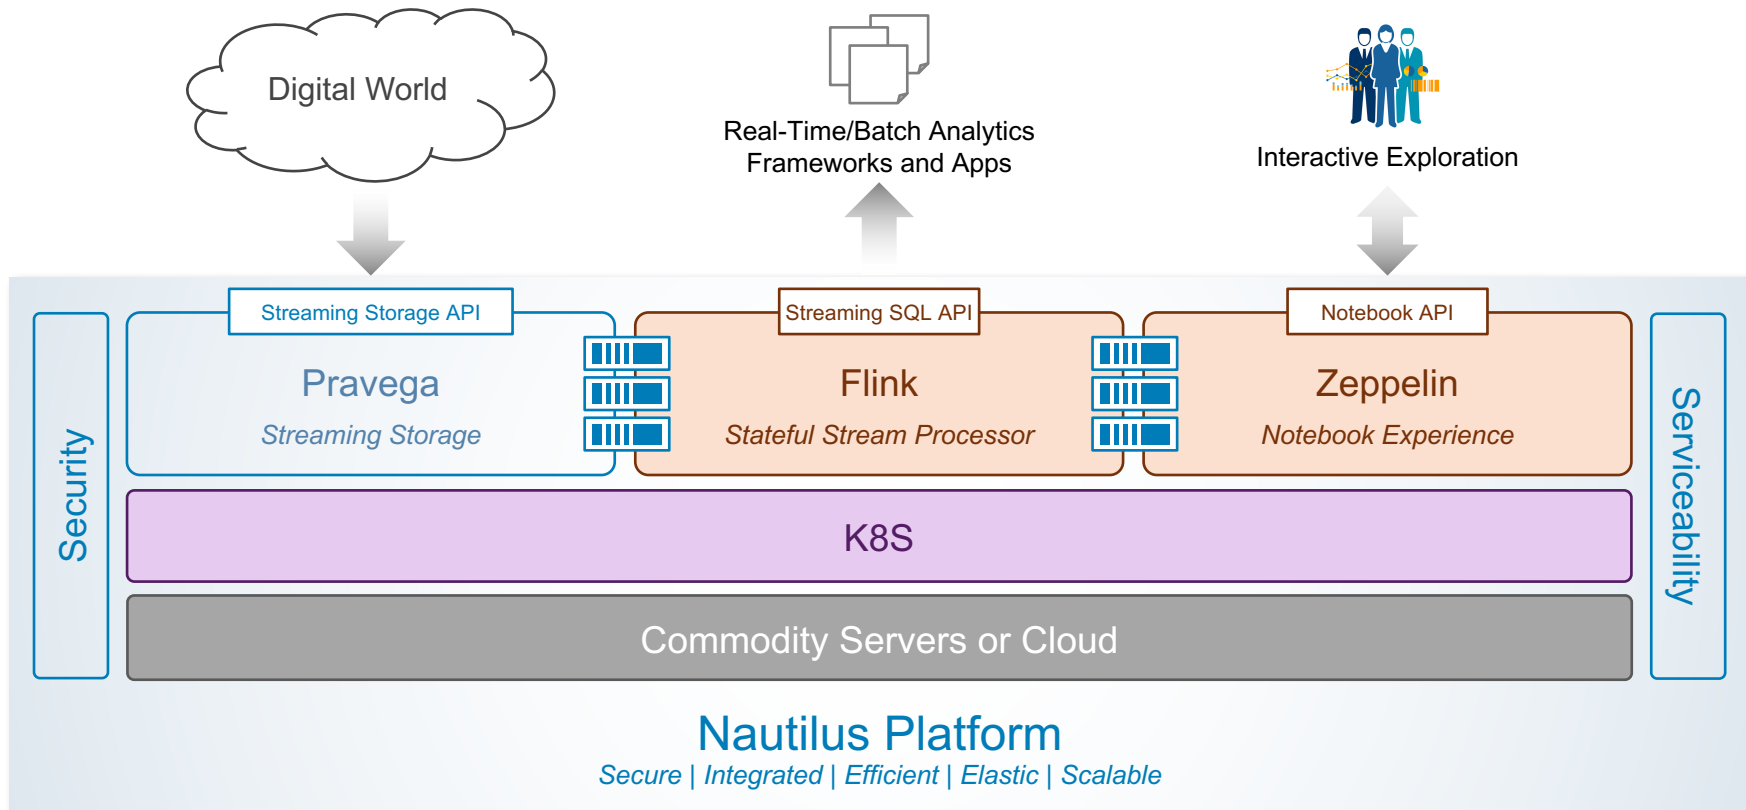
Transactional Semantics For “Exactly Once”



Pravega Optimizations for Stream Processors



A Turn-Key Streaming Data Platform



Summary

1. “Streaming Architecture” replaces “Accidental Architecture”
 - Data: infinite/continuous vs. static/finite
 - Correctness in real-time: Exactly once processing + consistent storage
2. Pravega Streaming Storage Enables Storage Refactoring
 - Infinite, durable, scalable, re-playable, elastic append-only log
 - Open source project
3. Unified Storage + Unified Data Pipeline
 - The New Data Stack!

D~~EL~~LEMC

CNUTCon [上海]
全球运维技术大会

主办方 Geekbang InfoQ
极客邦科技

50+ 年末充电[⚡]

开发&运维技术干货大盘点

容器

Kubernetes

DevOps

全链路压测

Severless

自动化运维

Service Mesh

Elasticsearch

微服务

使用折扣码 「QCon」 优惠报名 咨询电话：13269078023



扫码锁定席位

QCon

全球软件开发大会

北京·2019

更多技术干货分享，北京站精彩继续
提前参与，还能享受更多优惠

识别二维码
查看了解更多

2019.qconbeijing.com



Comparing Pravega and Kafka Design Points

Unlike Kafka, Pravega is designed to be a durable and permanent storage system

Quality	Pravega Goal	Kafka Design Point
Data Durability	Replicated and persisted to disk before ACK	Replicated but not persisted to disk before ACK ❌
Strict Ordering	Consistent ordering on tail and catch-up reads	Messages may get reordered ❌
Exactly Once	Producers can use transactions for atomicity	Messages may get duplicated ❌
Scale	Tens of millions of streams per cluster	Thousands of topics per cluster ❌
Elastic	Dynamic partitioning of streams based on load and SLO	Statically configured partitions ❌
Size	Log size is not bounded by the capacity of any single node	Partition size is bounded by capacity of filesystem on its hosting node ❌
	Transparently migrate/retrieve data from Tier 2 storage for older parts of the log	External ETL required to move data to Tier 2 storage; no access to data via Kafka once moved ❌
Performance	Low (<10ms) latency durable writes; throughput bounded by network bandwidth	Low-latency achieved only by reducing replication/reliability parameters ❌
	Read pattern (e.g. many catch-up readers) does not affect write performance	Read patterns adversely affects write performance due to reliance on OS filesystem cache ❌